

# The Particle-In-Cell method: lesson 2

**Elisabetta Boella**

e.boella@lancaster.ac.uk

<https://www.lancaster.ac.uk/staff/boella/>

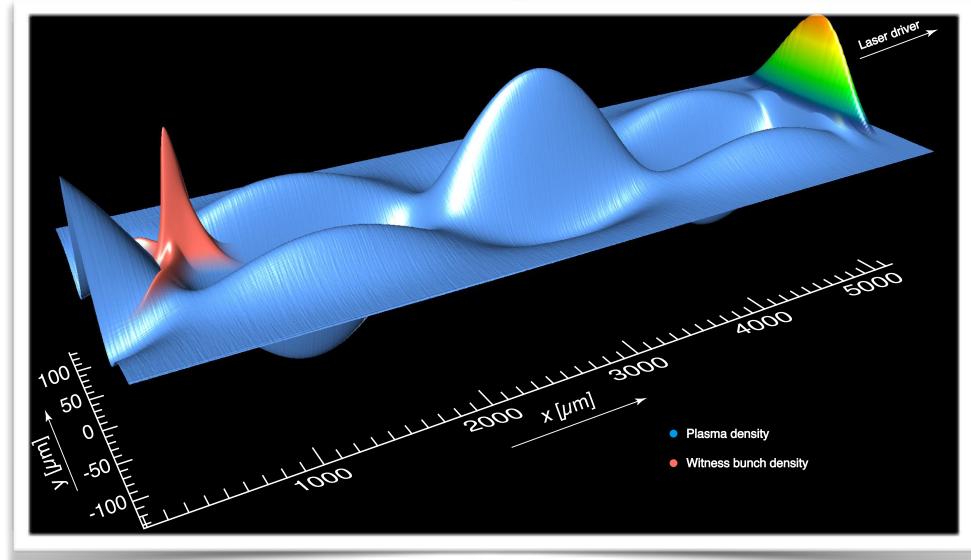


Physics Department

Lancaster University (Lancaster, UK)

Cockcroft Institute

Sci-Tech Daresbury (Warrington, UK)



## Acknowledgements

---

- ❖ Ricardo Fonseca (Instituto Superior Técnico, Lisbon, Portugal)
- ❖ Maria Elena Innocenti (KU Leuven, Leuven, Belgium)
- ❖ Giovanni Lapenta (KU Leuven, Leuven, Belgium)

# Higher order interpolation functions reduce numerical noise

## For simplicity 1D electrostatic case

Fourier transform of the force applied to one macroparticle:  $F(k) = \int_{-\infty}^{\infty} F(x)e^{-ikx} dx$

$$\text{Force applied to particle } i: F_i = q_i \Delta x \sum_j E_j W(x_j - x_i)$$

electric field @grid node j   interpolation function from grid node j to particle i

Inserting the expression for the force into the expression for the Fourier transform, we get

$$F(k) = q \left[ \Delta x \sum_j E_j e^{-ikx_j} \right] \int_{-\infty}^{\infty} W(x_j - x) e^{ik(x_j - x)} dx = q E(k) W(-k)$$

interpolation function  
Fourier transform

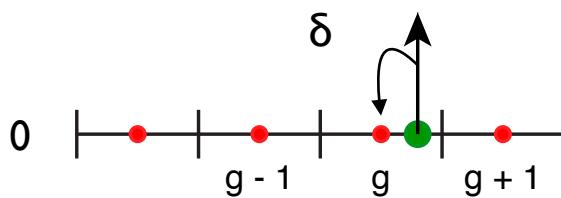
Discrete Fourier transform of  $E_j$   
Periodic

Importance of high k-modes  
depend of how fast  $W(k) \rightarrow 0$   
for  $k \rightarrow \pm \infty$

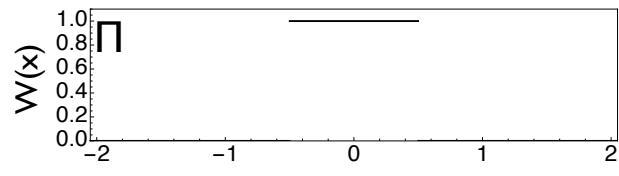
# Higher order interpolation functions reduce numerical noise

Order

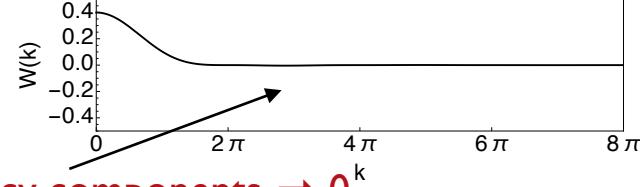
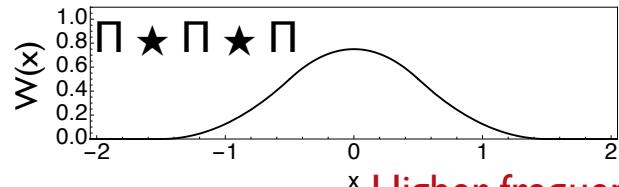
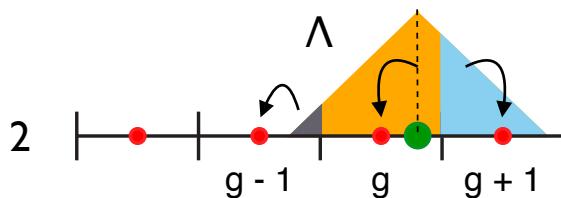
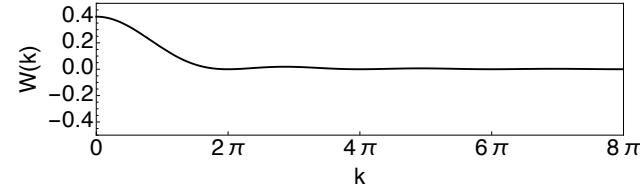
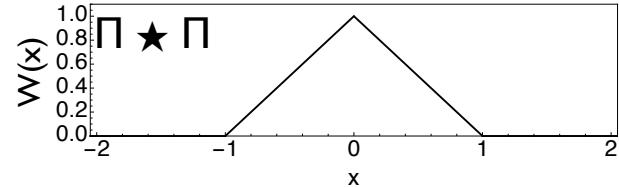
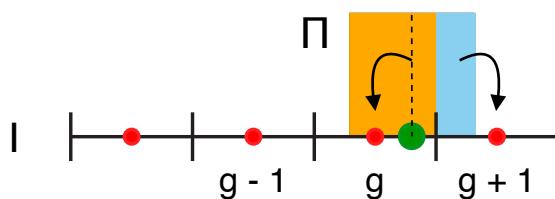
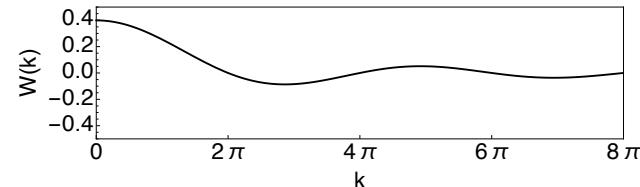
Shape function



Interpolation function



Fourier transform of  $W(x)$



Higher frequency components  $\rightarrow 0$

# Numerical heating due to aliasing can be reduced by resolving the Debye length

**For simplicity 1D electrostatic case**

Let us define the density of particle centres as  $n(x) = \sum_{i=1}^{N_p} \delta(x - x_i)$

Then the charge density on the grid will be  $\rho(x_g) = \frac{q}{\Delta x} \int n(x') W(x_g - x') dx'$

Similarly the charge density in the continuum will be  $\rho'(x) = \frac{q}{\Delta x} \int n(x') W(x - x') dx' = \frac{q}{\Delta x} n(x) \star W(x)$

with Fourier transform  $[\mathcal{F}\rho'(x)](k) = \hat{\rho}'(k) = \frac{q}{\Delta x} \hat{n}(k) \hat{W}(k)$

where the following definitions are valid  $\hat{\rho}'(k) = \int_{-\infty}^{\infty} \rho'(x) e^{-ikx} dx$  and  $\rho'(x) = \int_{-\infty}^{\infty} \frac{1}{2\pi} \hat{\rho}'(k) e^{ikx} dk$

## Numerical heating due to aliasing can be reduced by resolving the Debye length

The Fourier transform of the discrete charge density is  $[\mathcal{F}\rho(x_g)](k) = \hat{\rho}(k) = \Delta x \sum_{g=-\infty}^{\infty} \rho(x_g) e^{-ikx_g}$

with  $\rho(x_g) = \int_{-\pi/\Delta x}^{\pi/\Delta x} \frac{1}{2\pi} \hat{\rho}(k) e^{ikx_g} dk \quad (1)$

At  $x_g \rho(x_g) = \rho'(x_g) \Rightarrow \rho(x_g) = \int_{-\infty}^{\infty} \frac{1}{2\pi} \hat{\rho}'(k) e^{ikx_g} dk = \int_{-\pi/\Delta x}^{\pi/\Delta x} \frac{1}{2\pi} \sum_{n=-\infty}^{\infty} \hat{\rho}'(k-nk_g) e^{ikx_g} dk' \quad (2)$

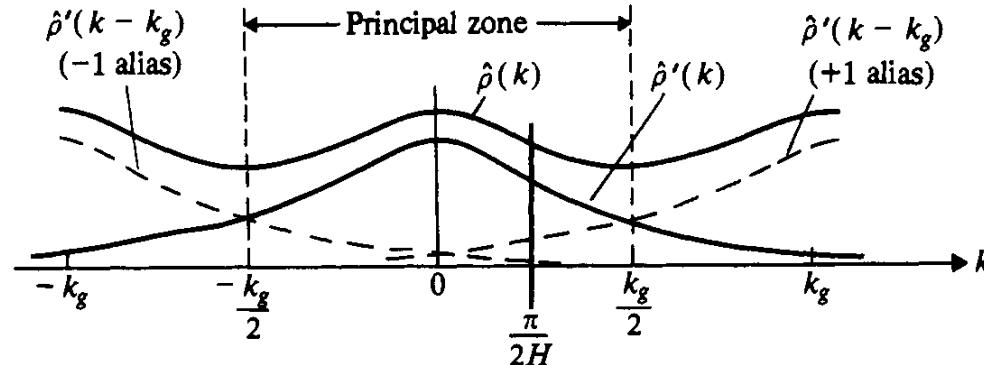
splitting the integral into segments of length  $k_g = 2\pi/\Delta x$

changing variable  $k' = k - nk_g$

Comparing (1) and (2) we get  $\hat{\rho}(k) = \sum_{n=-\infty}^{\infty} \hat{\rho}'(k - nk_g) = \frac{q}{\Delta x} \sum_{n=-\infty}^{\infty} \hat{n}(k - nk_g) \hat{W}(k - nk_g)$

The Fourier transform of the discrete charge density is given by the sum of all the replicas of the Fourier transform of the continuous charge density. This gives rise to a phenomena called aliasing for not properly selected values  $k_g = 2\pi/\Delta x$ .

# Numerical heating due to aliasing can be reduced by resolving the Debye length



Aliasing derives from the fact that we are trying to squeeze continuous information into a discrete grid

The aliasing effects in k-space is fed back into the simulation producing numerical noise in the real space

Things become worse and worse if we do not sample enough, e.g. if we consider  $\Delta x$  bigger and bigger  
 Things become better with higher order interpolation functions whose Fourier transform is 0 for high  $k_s$

Aliasing becomes a serious problem if  $\Delta x \gg \lambda_D$  causing the development of the finite grid instability, which results in unphysical electron heating. In the case of linear interpolation, the instability is avoided if  $\Delta x/\lambda_D \ll \pi$

# The described PIC algorithm conserves total momentum

**For simplicity 1D electrostatic case**

$$\frac{dP}{dt} = \sum_i F_i = \sum_i q_i \Delta x \sum_j E_j W(x_j - x_i) = \Delta x \sum_j E_j \sum_i q_i W(x_j - x_i) = \Delta x \sum_j \rho_j E_j$$

↓  
 electric field interpolated  
 at particle position  $x_i$ 

 ↓  
 density on the grid  
 at  $x_j$

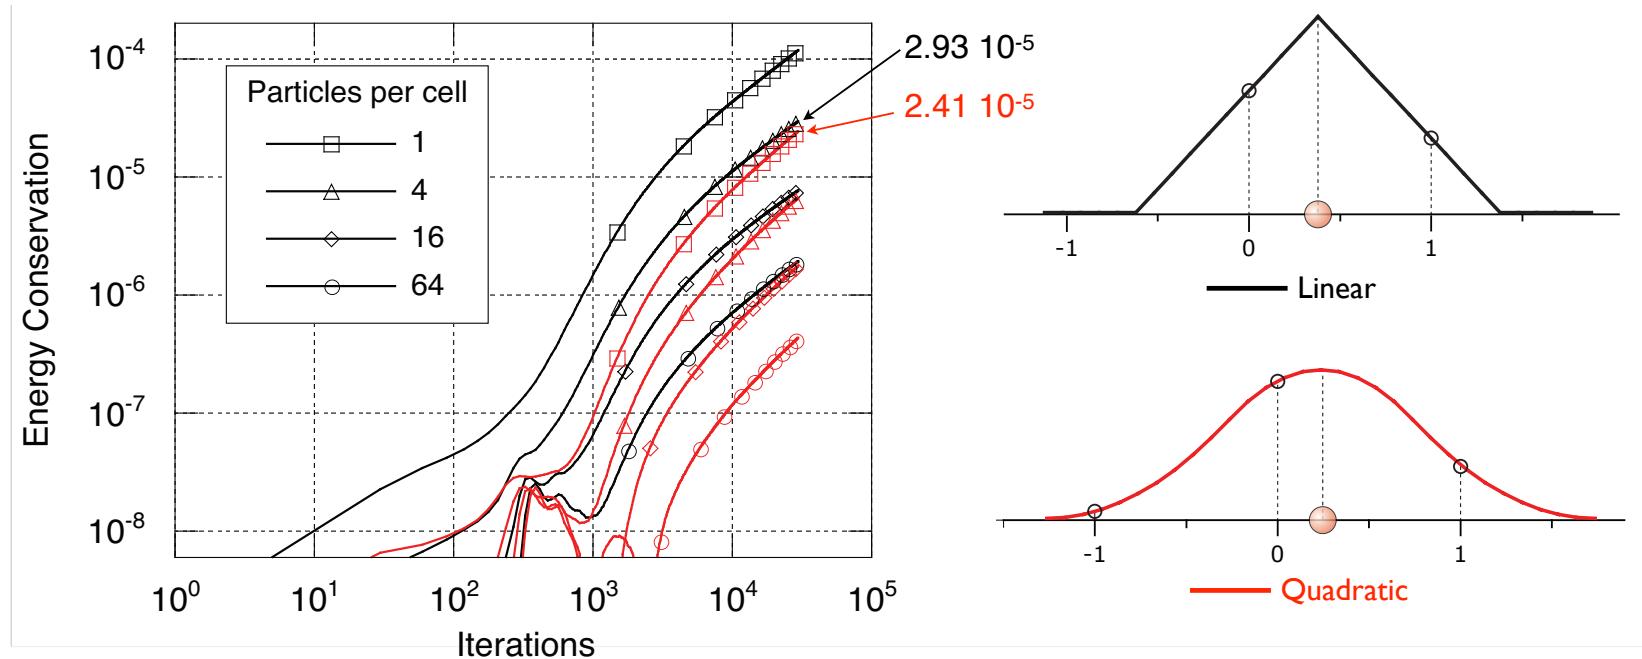
In a periodic system with left-right symmetry  $\Delta x \sum_j \rho_j E_j = 0$

Therefore momentum is conserved!

When momentum is conserved, the total energy is not conserved.

Therefore, it is important to always monitor the evolution of the total energy of your simulation.

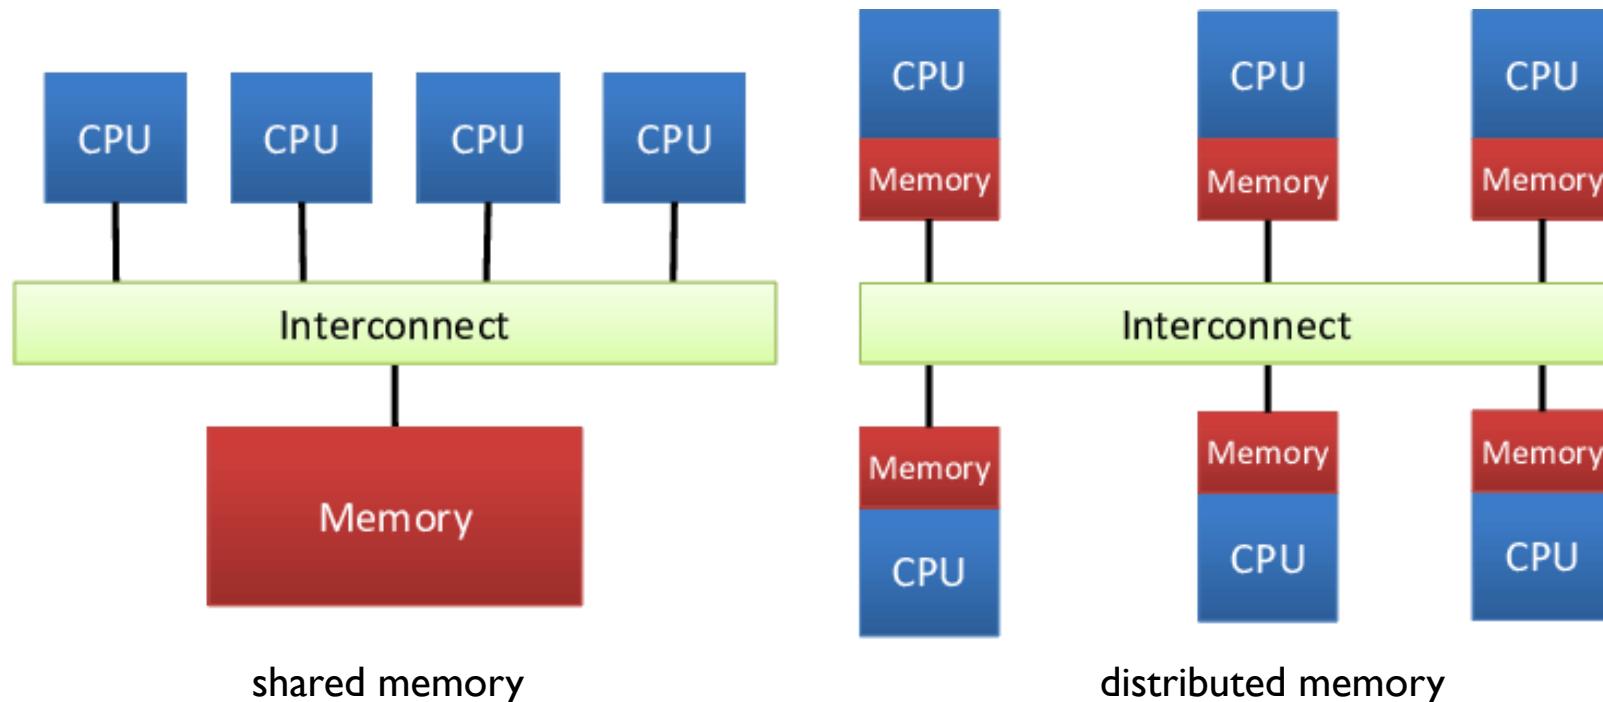
# Higher order interpolation and many particle per cell improve energy conservation



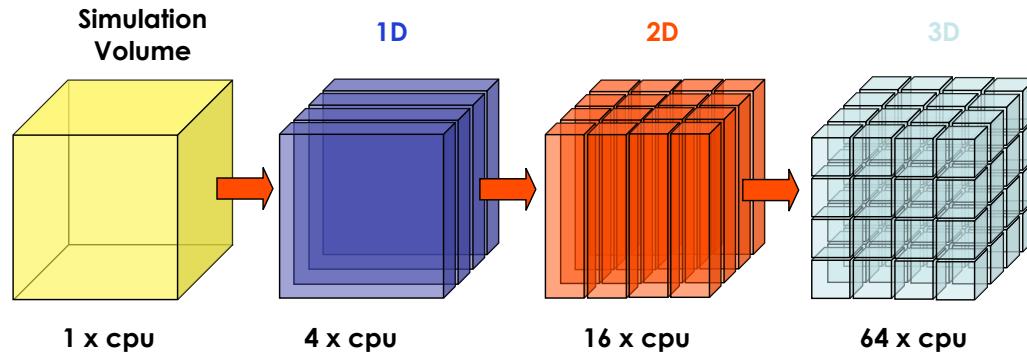
# Realistic PIC simulations need supercomputers



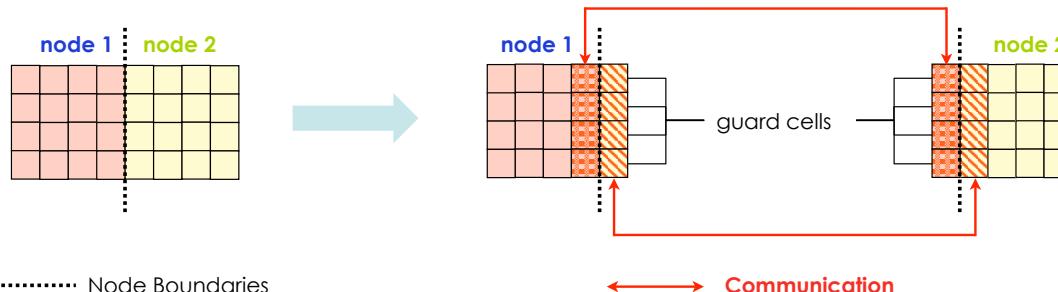
## Different parallelism paradigms



# Domain decomposition: simulation space is shared among several CPUs



Guard cells are used to communicate among the processors



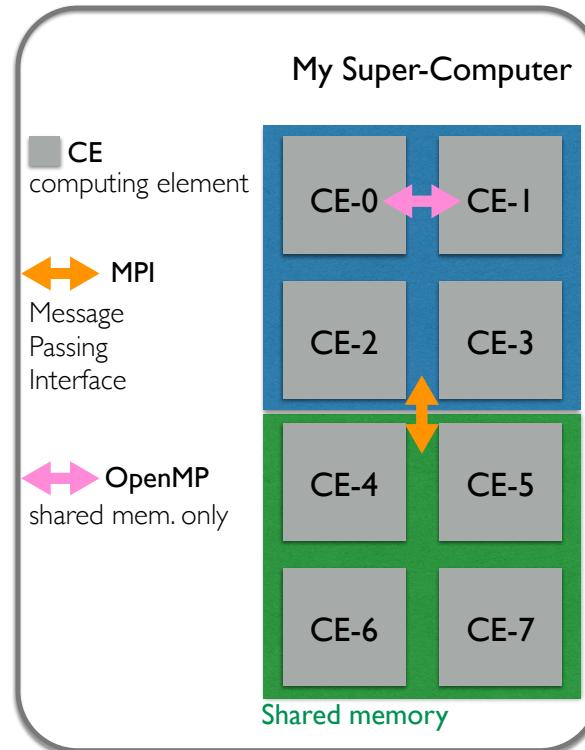
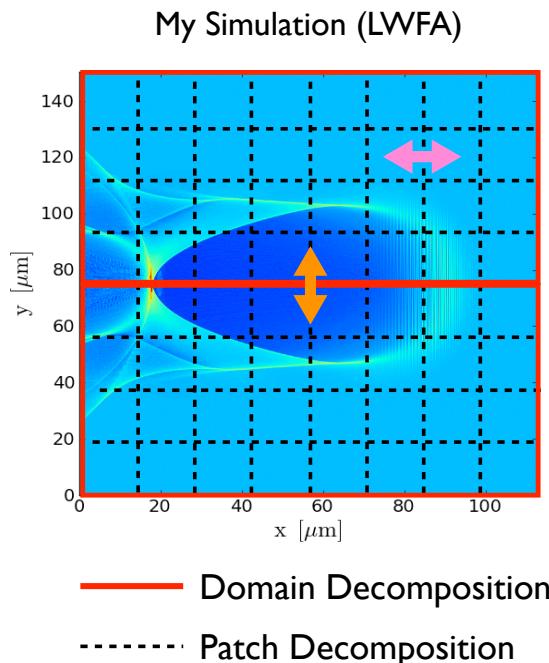
**Usually field solver is local**

Communication only with neighbouring nodes

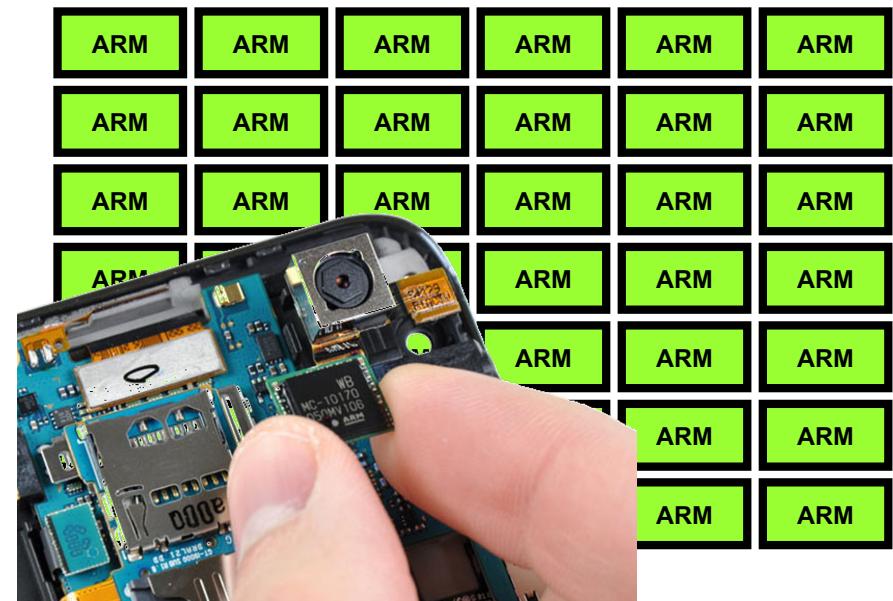
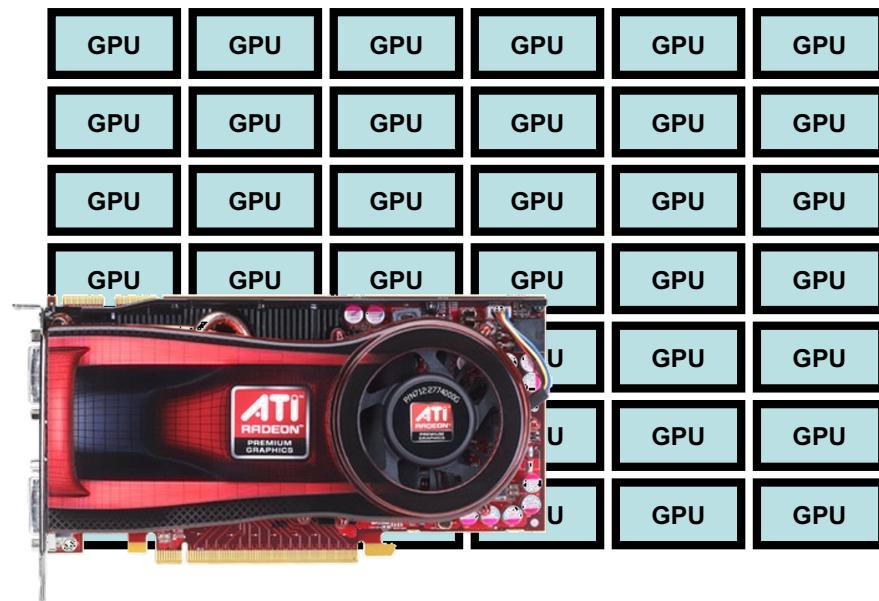
**Usually the same decomposition for particles and fields is used**

Each computing node acts as a section of the global physical space

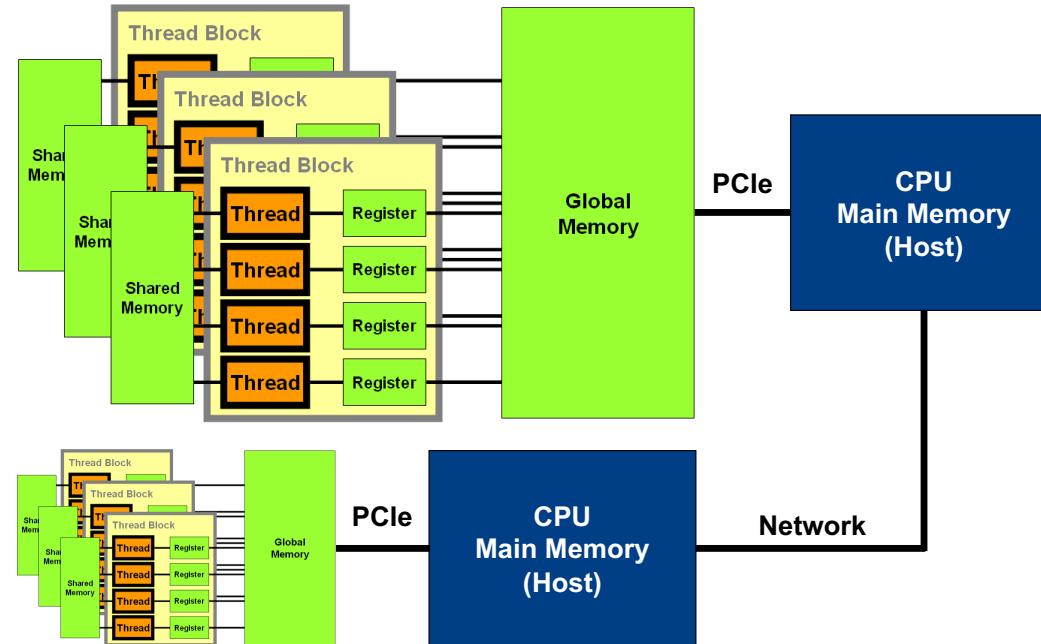
# Hybrid parallelism allows for better taking advantage of HPC resources



# Towards the exascale



# Memory management becomes even more critical: it is important to keep data locality



Structure of arrays vs **array of structures**

## Testing the performances: strong scaling

### Scaling of the code PIConGPU on the supercomputer D.A.V.I.D.E. @CINECA (Italy)

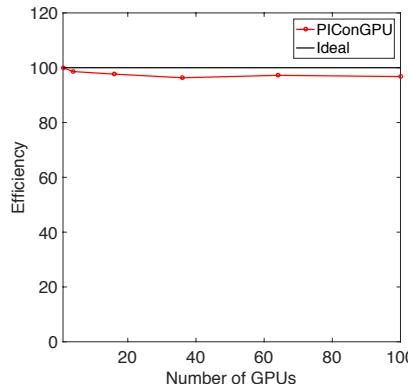
Grid size: 1200x1200

Super cell size: 16x16

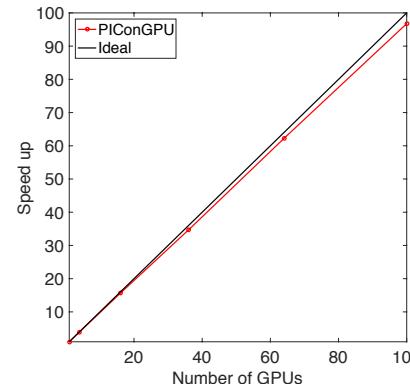
Time steps: 1000

Particles/cell: 8x8

Interpolation: linear



Efficiency @10x10 GPUs: 96.77%



Speed up @10x10 GPUs: 96.77

## Testing the performances: weak scaling

### Scaling of the code PIConGPU on the supercomputer D.A.V.I.D.E. @CINECA (Italy)

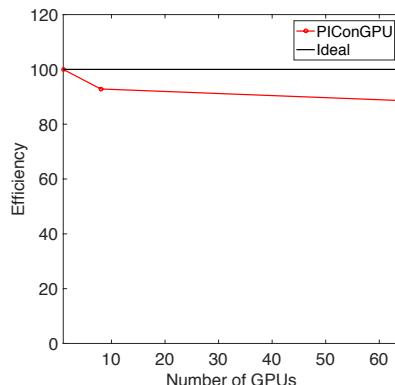
Grid size:  $64 \times 64 \times 64 \rightarrow 320 \times 320 \times 320$

Super cell size:  $8 \times 8 \times 4$

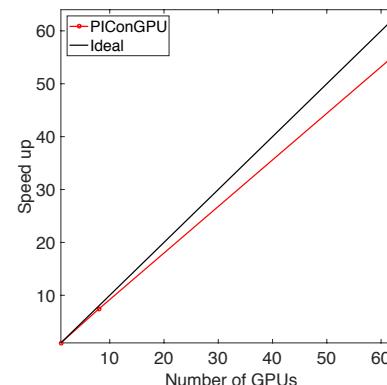
Time steps: 1000

Particles/cell:  $4 \times 2 \times 2$

Interpolation: quadratic

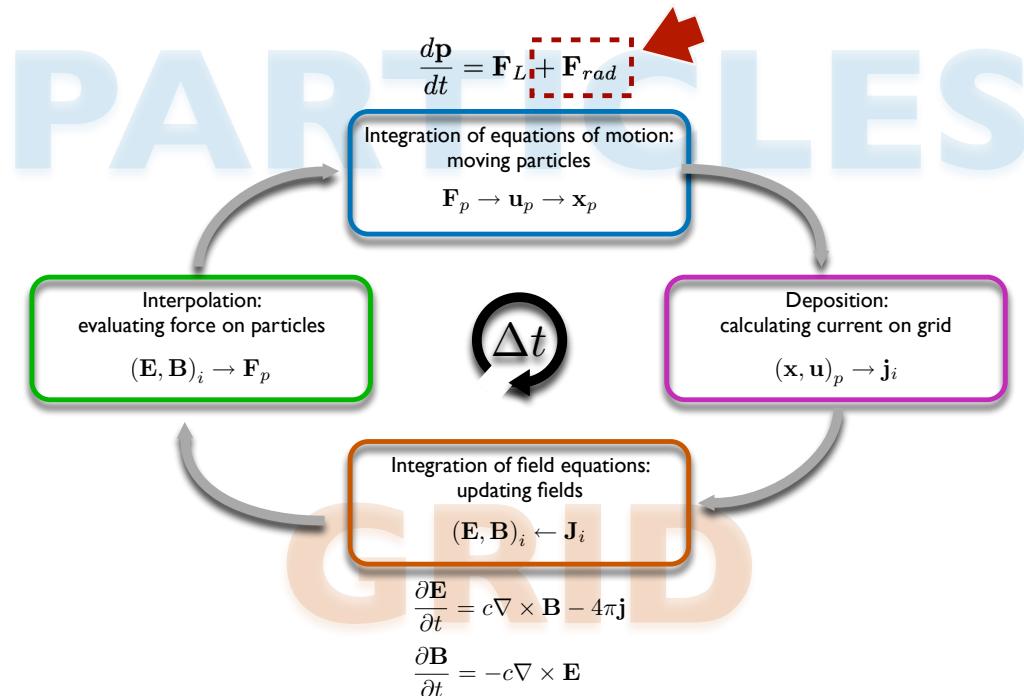


Efficiency @ $4 \times 4 \times 4$  GPUs: 88.64%



Speed up @ $4 \times 4 \times 4$  GPUs: 56.73

# Beyond the classical PIC algorithm: modelling classical and quantum radiation reaction



# Beyond the classical PIC algorithm: modelling classical and quantum radiation reaction

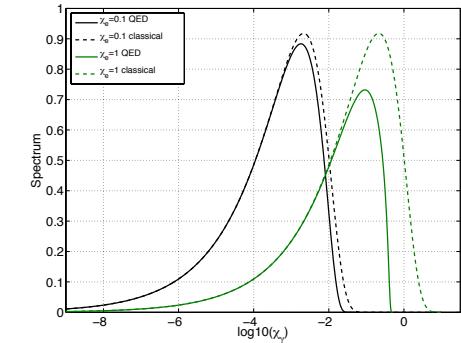
## Different types of Radiation reaction models

$$\frac{dp}{dt} = \mathbf{F}_L + \begin{cases} \mathbf{F}_{rad} & \text{Continuous damping rate} \\ \frac{d^2 P}{dtd\chi_\gamma} & \text{QED probabilistic approach} \end{cases}$$

## Implementation in PIC codes

- ❖ Continuous damping rate: particle pusher with  $\mathbf{F}_{rad}$
- ❖ QED probabilistic approach: particle pusher + Monte Carlo module  
 every  $\Delta t$  : probability of photon emission  
 Select a photon in QED synchrotron spectrum  
 Update particle momentum due to quantum recoil
- ❖ The QED approach can be generalized to any external EM fields under the conditions:
  - quasi-static fields  $t_{carac}(\vec{E}, \vec{B}) \gg t_{coh} \implies \gamma \gg 1$
  - weak fields  $\chi_e^2 \gg \text{Max}(f, g) \quad (f, g) \ll 1$

Synchrotron Spectrum



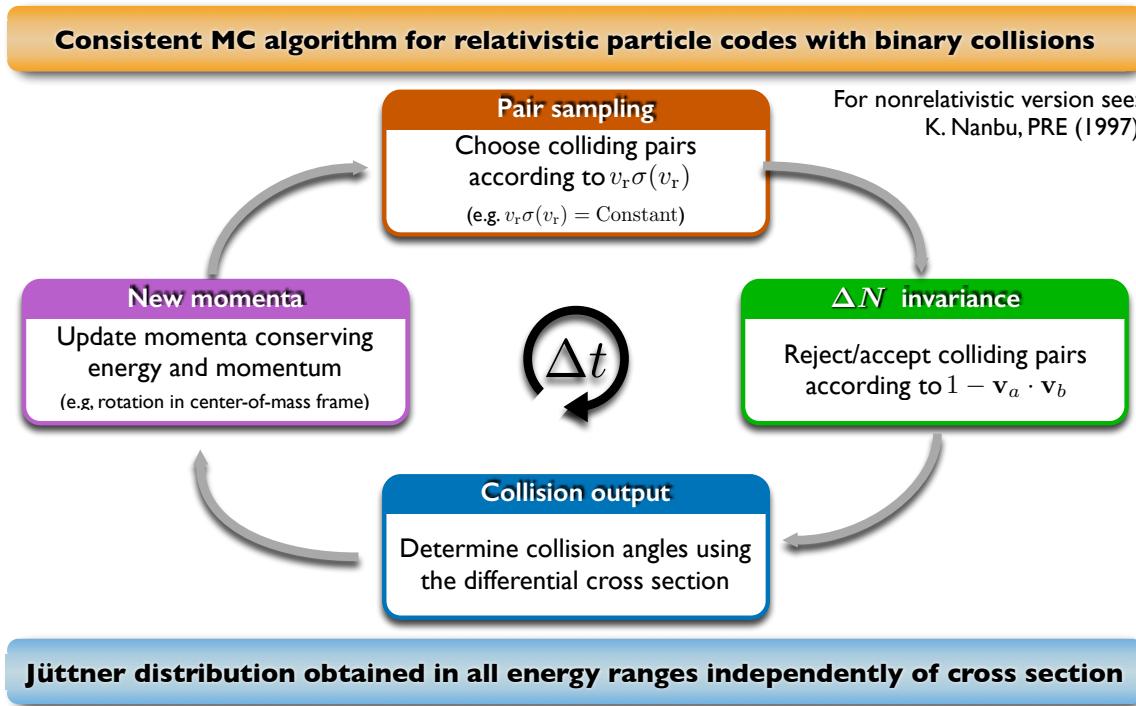
$$f = F_{\mu\nu}^2 / E_{crit}^2$$

$$g = F_{\mu\nu}^* F_{\mu\nu} / E_{crit}^2$$

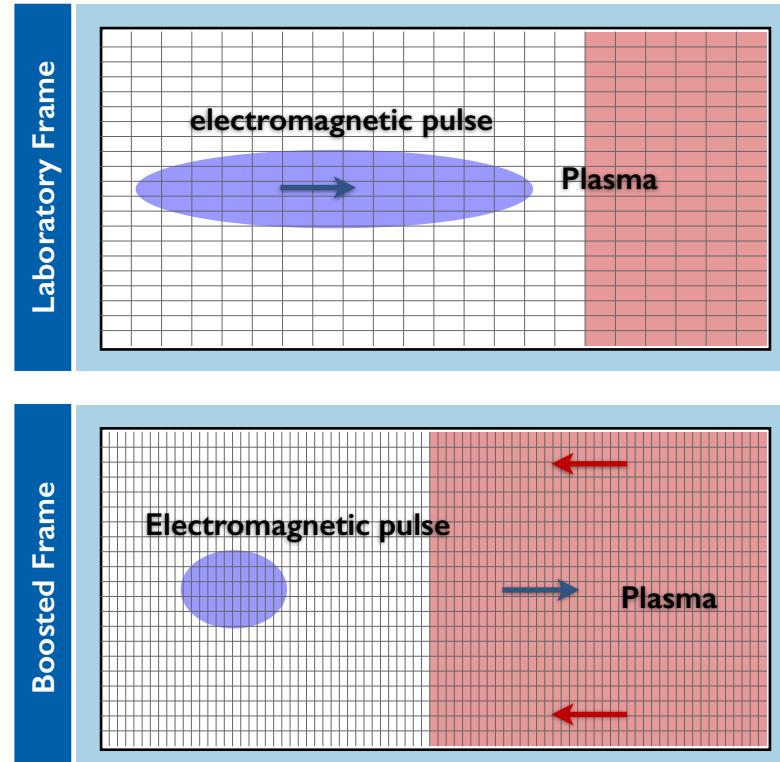
$$E_{crit} = m^2 c^3 / e\hbar$$

$$\chi_{e,\gamma} = \frac{|F_{\mu\nu} p_{e,\gamma}^\nu|}{E_{crit} mc}$$

# Beyond the classical PIC algorithm: modelling Coulomb collision



# Beyond the classical PIC algorithm: modelling LWFA in a boosted frame



## Hands-on: electromagnetic PIC

---

Download the zPIC code from the GitHub repo

git clone <https://github.com/ricardo-fonseca/zpic.git>

or download from <https://github.com/ricardo-fonseca/zpic>

Go to the folder python and compile the code using make

Launch the Jupyter notebook from the source folder:

> jupyter notebook LWFA1D.ipynb

We are going to model plasma-based accelerators

# Hands-on: electromagnetic PIC

## Choose the normalisation

### Plasma sets reference

Plasma density is unity

- ❖ Normalise lengths to plasma skin depth and frequency to plasma frequency

#### Example

- ❖ Plasma density  $n_p = 10^{18} \text{ cm}^{-3}$
- ❖ Plasma frequency  $\omega_p \sim 5.64 \times 10^{13} \text{ rad s}^{-1}$
- ❖ Laser wavelength  $\lambda_0 = 1 \mu\text{m}$
- ❖ Laser frequency  $\omega_0 \sim 2.34 \times 10^{15} \text{ rad s}^{-1}$
- ❖ Normalised laser frequency is  $\omega_0/\omega_p \sim 41.5$

### Laser sets reference

Reference laser frequency is unity

- ❖ Normalise plasma density to critical density; length to inverse laser wavenumber

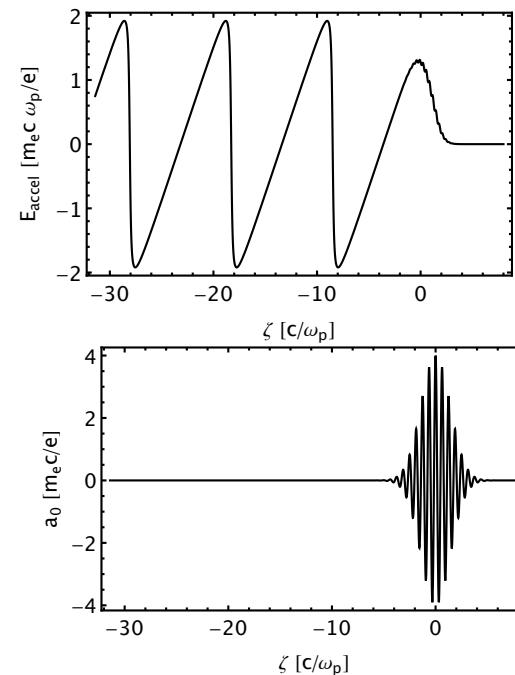
#### Example

- ❖ Laser wavelength  $\lambda_0 = 1 \mu\text{m}$
- ❖ Laser frequency  $\omega_0 \sim 2.34 \times 10^{15} \text{ rad s}^{-1}$
- ❖ Critical frequency  $n_{\text{crit}} \sim 1.72 \times 10^{21} \text{ cm}^{-3}$
- ❖ Plasma density  $n_p = 10^{18} \text{ cm}^{-3}$
- ❖ Normalised plasma density  $n_p/n_{\text{crit}} \sim 5.8 \times 10^{-4}$

# Hands-on: electromagnetic PIC

## LWFA simulation

- ❖ Laser propagates in an underdense plasma  
 $n_p \ll n_{crit} \mid \lambda_0 \ll \lambda_p \mid \omega_p \ll \omega_0$
- ❖ Need to resolve the smallest scale length  
 20 - 30 cells per wavelength
- ❖ Plasma wave  
 Skin depth sets the plasma scale length  
 $c/\omega_p \sim 5.3 \text{ } \mu\text{m}/( n_p [10^{18} \text{ cm}^{-3}] )^{1/2}$
- ❖ Laser  
 laser wavelength sets the laser scale length  
 $\lambda_0 \sim 1 \text{ } \mu\text{m} \sim 0.18 ( n_p [10^{18} \text{ cm}^{-3}] )^{1/2} c/\omega_p$

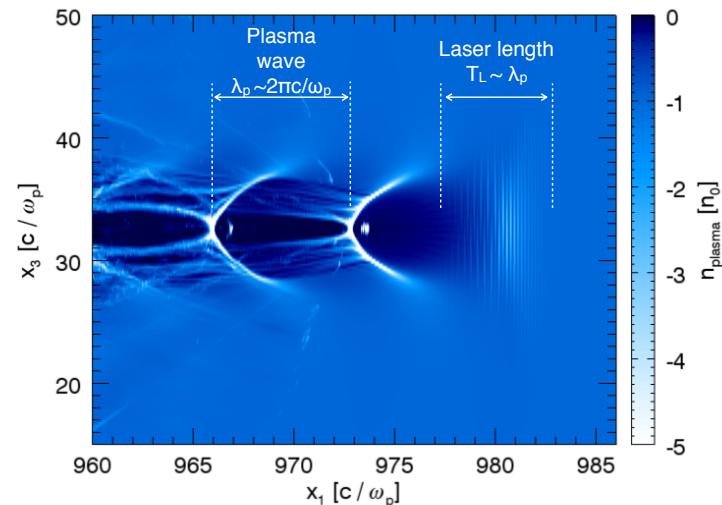


Longitudinal spatial Resolution:  $\Delta x \sim \lambda_0 / \# \sim 0.18 / \# ( n_p [10^{18} \text{ cm}^{-3}] )^{1/2} c/\omega_p$   
 $\# > 20-30$  (number of cells per laser wavelength)

# Hands-on: electromagnetic PIC

## Determine simulation box size

- ❖ Simulations are done in a moving window moving at the speed of light  
The simulation box does not need to hold the entire propagation length
- ❖ Simulation box needs only to model the relevant structures in the accelerator  
Laser driver and initial trailing buckets of accelerating structure
- ❖ Box size determined by largest relevant structures  
Longitudinally
  - a few plasma wavelengths long
  - $> 4 \lambda_p \sim 25 c/\omega_p$
 Transversely (2D)
  - Laser pulse waist / transverse bubble size
  - $> 4 \lambda_p \sim 25 c/\omega_p$





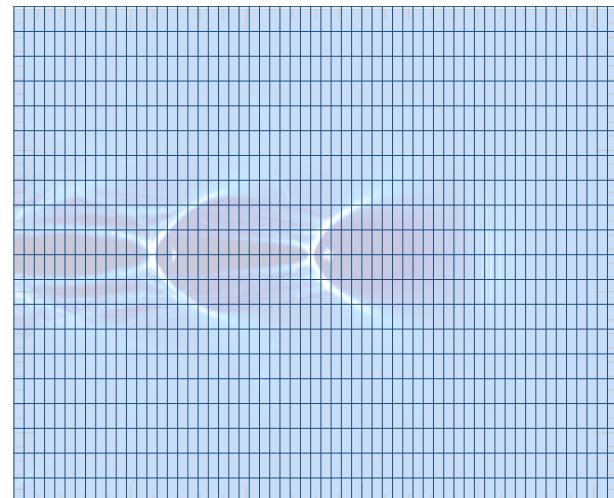
# Hands-on: electromagnetic PIC

## Set simulation grid

- ❖ Box length:  $L = 4 \lambda_p$
- ❖ 20 points per laser wavelength
- ❖  $\Delta x \sim \lambda_0/20 \sim 0.18/20 = 0.009 \text{ c}/\omega_p$
- ❖ Number of cells  $\sim L / \Delta x \sim 2800$  cells

## Choose simulation particles

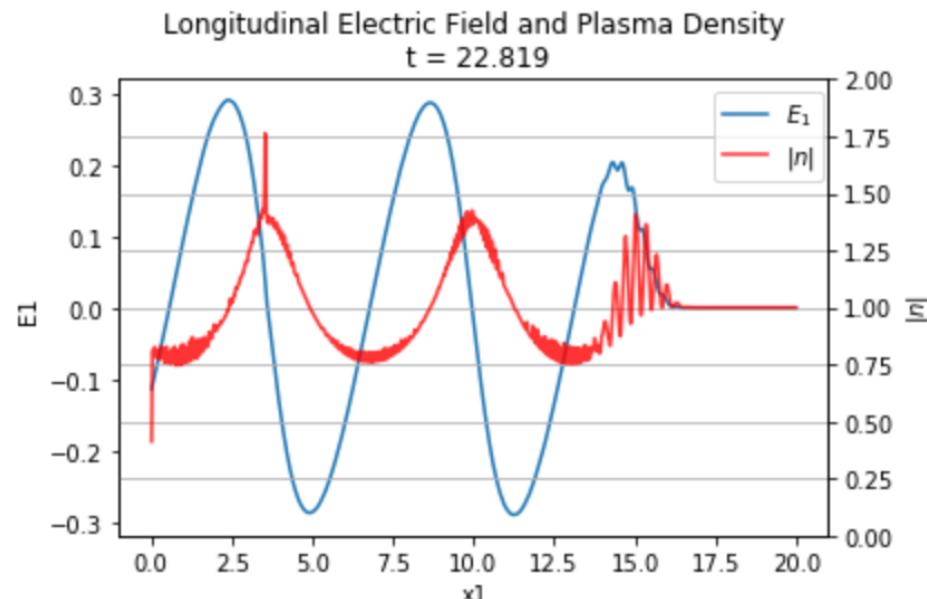
- ❖ Number of particles per cell must resolve local phasespace
- ❖  $\gg 1$  in 1D (e.g. 64)
- ❖  $\sim 10$  in 2D
- ❖ Higher numbers improve phasespace resolution  
(detailed distribution tails) and reduces simulation noise



# Hands-on: electromagnetic PIC

## Choose diagnostic

- ❖ Plasma density  
Charge density of the background plasma  
Wave structure and particle loading
- ❖ Longitudinal electric field  
Accelerating / decelerating fields  
Transverse electric field  
Laser field  
Focusing / defocusing fields (2D)
- ❖ Particle phasespace  
Show particle momenta as a function of position  
Most common is  $u_1/x_1$   
Wave structure and particle acceleration



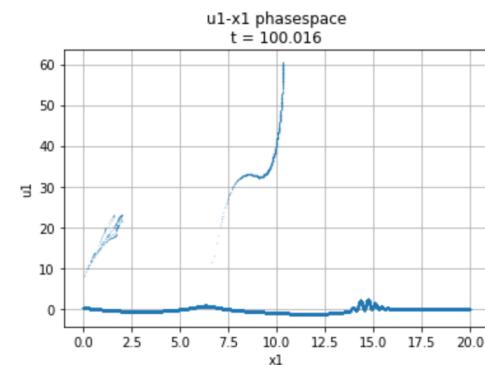
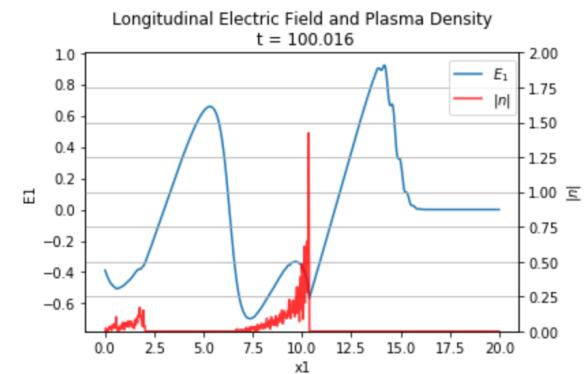
# Hands-on: electromagnetic PIC

## **Simulate a laser wakefield accelerator:**

- Add an ultra-intense laser beam as a driver ( $a_0 \sim 2$ )
- Choose laser length smaller than  $\lambda_p$

## **Questions:**

1. can you observe particle injection and trapping?
2. is the energy gain consistent with the longitudinal electric field values?
3. describe and justify the shape for the plasma electric field in the region where particles accelerate
4. could you accelerate positrons in this plasma wave? where would you place them and with what initial velocity/energy? try to simulate!



# Hands-on: electromagnetic PIC

## Laser Wakefield Accelerator

1D simulation of a laser wakefield accelerator.

```
In [1]: import emld
import numpy

# Time step
dt = 0.019

# Simulation time
tmax = 22.8

# Number of cells
nx = 1000

# Simulation box size
box = 20.0

## Background plasma

# Particles per cell
ppc = 128

# Use a step density profile
electrons = emld.Species( "electrons", -1.0, ppc,
                           density = emld.Density( type = "step", start = 20.0))

# Initialize simulation
sim = emld.Simulation( nx, box, dt, species = electrons )

# Add laser pulse
sim.add_laser( emld.Laser( start = 17.0, fwhm = 2.0, a0 = 1.0, omega0 = 10.0, polarization = numpy.pi/2 ) )

# Set moving window
sim.set_moving_window()

# Set current smoothing
sim.set_smooth( emld.Smooth(xtype = "compensated", xlevel = 4) )

# Run the simulation
sim.run( tmax )
```

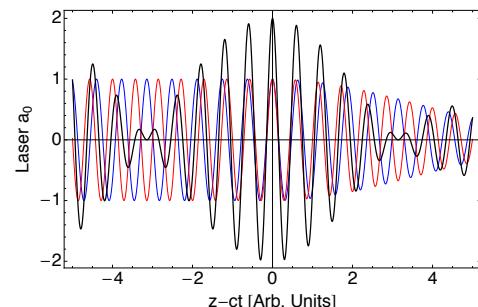
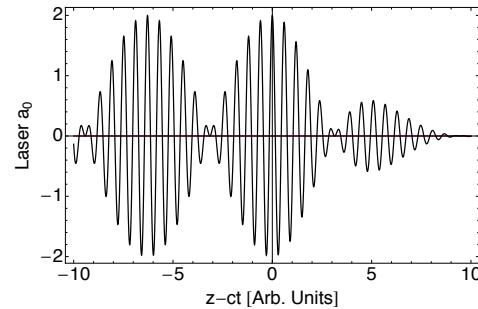
# Hands-on: electromagnetic PIC

## Simulate a plasma beat-wave accelerator:

- Super-impose three laser modes with frequencies differing by  $\omega_p$  (e.g.  $\omega_0 = 10, 11 \omega_p$ )
- Choose Laser length  $\gg \lambda_p$

## Questions:

1. why does the plasma wave amplitude increase along the pulse?
2. what happens if the initial frequencies of the lasers are not separated by the plasma frequency? Why?
3. compare the trapping threshold, as a function of the peak laser  $a_0$ , for a standard LWFA (with pulse length smaller than  $\lambda_p$ ) with the beat-wave accelerator. Which one is the lowest?
4. Decrease the amplitude of the laser side bands and run the simulation for longer times. What happens to the laser?



# Hands-on: electromagnetic PIC

## Laser Plasma Beatwave Accelerator

1D simulation of a laser plasma beatwave accelerator.

```
In [1]: import emld
import numpy

# Time step
dt = 0.019
# Simulation time
tmax = 100.
# Number of cells
nx = 1250
# Simulation box size
box = 50.0
# Particles per cell
ppc = 128

# Use a step density profile
electrons = emld.Species( "electrons", -1.0, ppc,
                           density = emld.Density( type = "step", start = 50.0))

# Initialize simulation
sim = emld.Simulation( nx, box, dt, species = electrons )

# Add laser pulses
sim.add_laser( emld.Laser( start = 50.0, rise = 20.0 , flat = 40.0 , fall = 1.0 , a0 = 0.5, omega0 = 10.0, polarization = 0.0 ) )
sim.add_laser( emld.Laser( start = 50.0, rise = 20.0 , flat = 40.0 , fall = 1.0 , a0 = 0.5, omega0 = 11.0, polarization = 1.0 ) )

# Set moving window
sim.set_moving_window()
# Set current smoothing
sim.set_smooth( emld.Smooth(xtype = "compensated", xlevel = 4) )

# Run the simulation
sim.run( tmax )
```

```
Running simulation up to t = 100 ...
n = 5264, t = 100.016
Done.
```

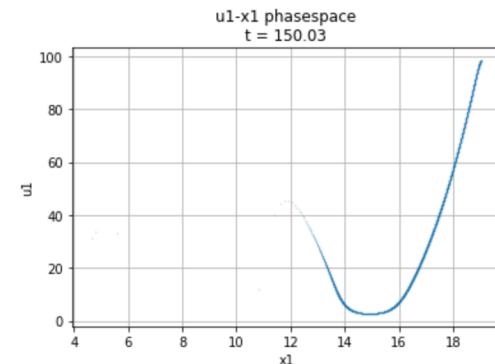
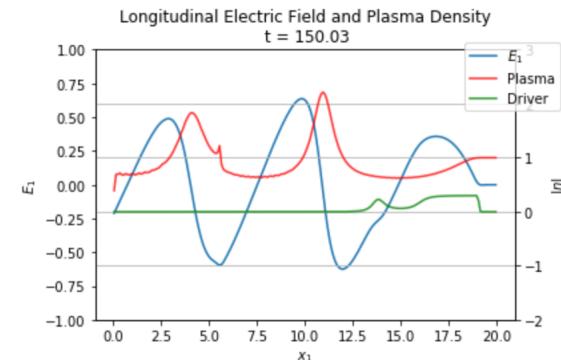
# Hands-on: electromagnetic PIC

## Simulate a plasma wakefield accelerator:

- Use an ultra-relativistic particle beam as a driver
  - e.g.  $u_{fl} = 100$ , length  $\sim 10 c/\omega_p$ , density  $\sim 0.3$
- Choose plasma length  $\gg \lambda_p$

## Questions:

1. Why does the head of the driver loose energy?
2. What happens to the energy of the driver if it has a length comparable to  $\lambda_p$ ?
3. What is the phase velocity of the plasma wave?
4. Can you observe plasma electron trapping and acceleration?



# Hands-on: electromagnetic PIC

## Plasma Wakefield Accelerator

1D simulation of a plasma wakefield accelerator.

```
In [63]: import emld
import numpy

# Time step
dt = 0.09
# Simulation time
tmax = 150.
# Number of cells
nx = 200
# Simulation box size
box = 20.0

# Particles per cell
ppc = 128

# Use a step density profile
electrons = emld.Species( "electrons", -1.0, ppc,
                           density = emld.Density( type = "step", start = 100.0))

driver = emld.Species( "driver", -1.0, ppc,
                       density = emld.Density( n = 0.3 , type = "slab", start = 8.0 , end = 19.0 ),
                       ufl=[100.0,0.0,0.0])

# Initialize simulation
sim = emld.Simulation( nx, box, dt, species = [electrons,driver] )

# Set moving window
sim.set_moving_window()

# Set current smoothing
sim.set_smooth( emld.Smooth(xtype = "compensated", xlevel = 4) )

# Run the simulation
sim.run( tmax )

Running simulation up to t = 150 ...
n = 1667, t = 150.03
Done.
```

# A beam means violating charge and current neutrality: it requires computing the initial electromagnetic fields

## Fields are calculated in the beam reference frame

Deposit beam charge in simulation frame

Boost to beam frame

Calculate electric field (no magnetic field in this frame)

Boost back to simulation frame, calculating magnetic field

## Electric field calculations

Are performed using Coulomb's law • This assumes open boundaries

In the beam frame each charge cell becomes elongated along propagation direction

Assume each cell is an infinite slab (2D) or rod (3D) • Only transverse fields!

